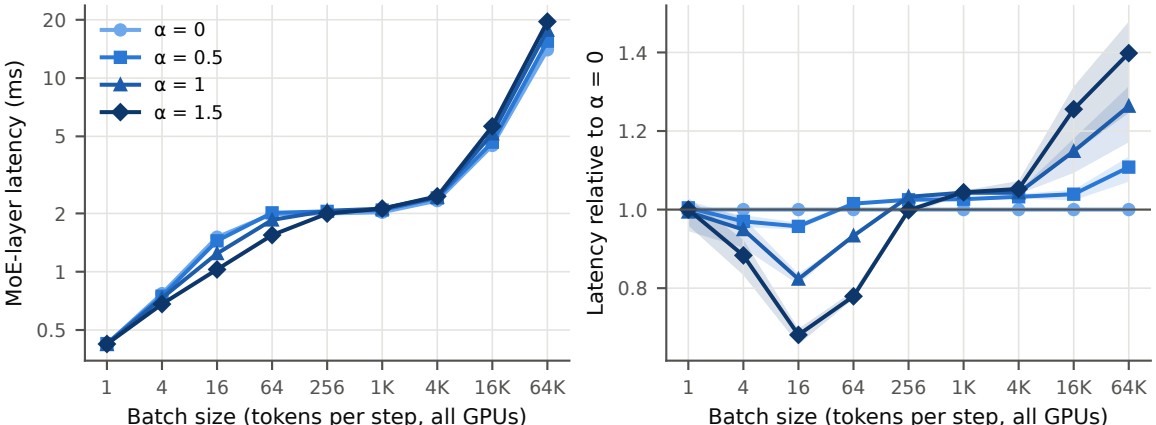
(a) Latency, 2 GPUs, No Load Balance
MoE-layer latency (ms)
20
10
5
2
1
0.5
α = 0
α = 0.5
α = 1
α = 1.5
1
4
16
64
256
1K
4K
16K
64K
Batch size (tokens per step, all GPUs)
(b) Effect of skew (band: range over 3 seeds)
Latency relative to α = 0
1.4
1.2
1.0
0.8
1
4
16
64
256
1K
4K
16K
64K
Batch size (tokens per step, all GPUs)