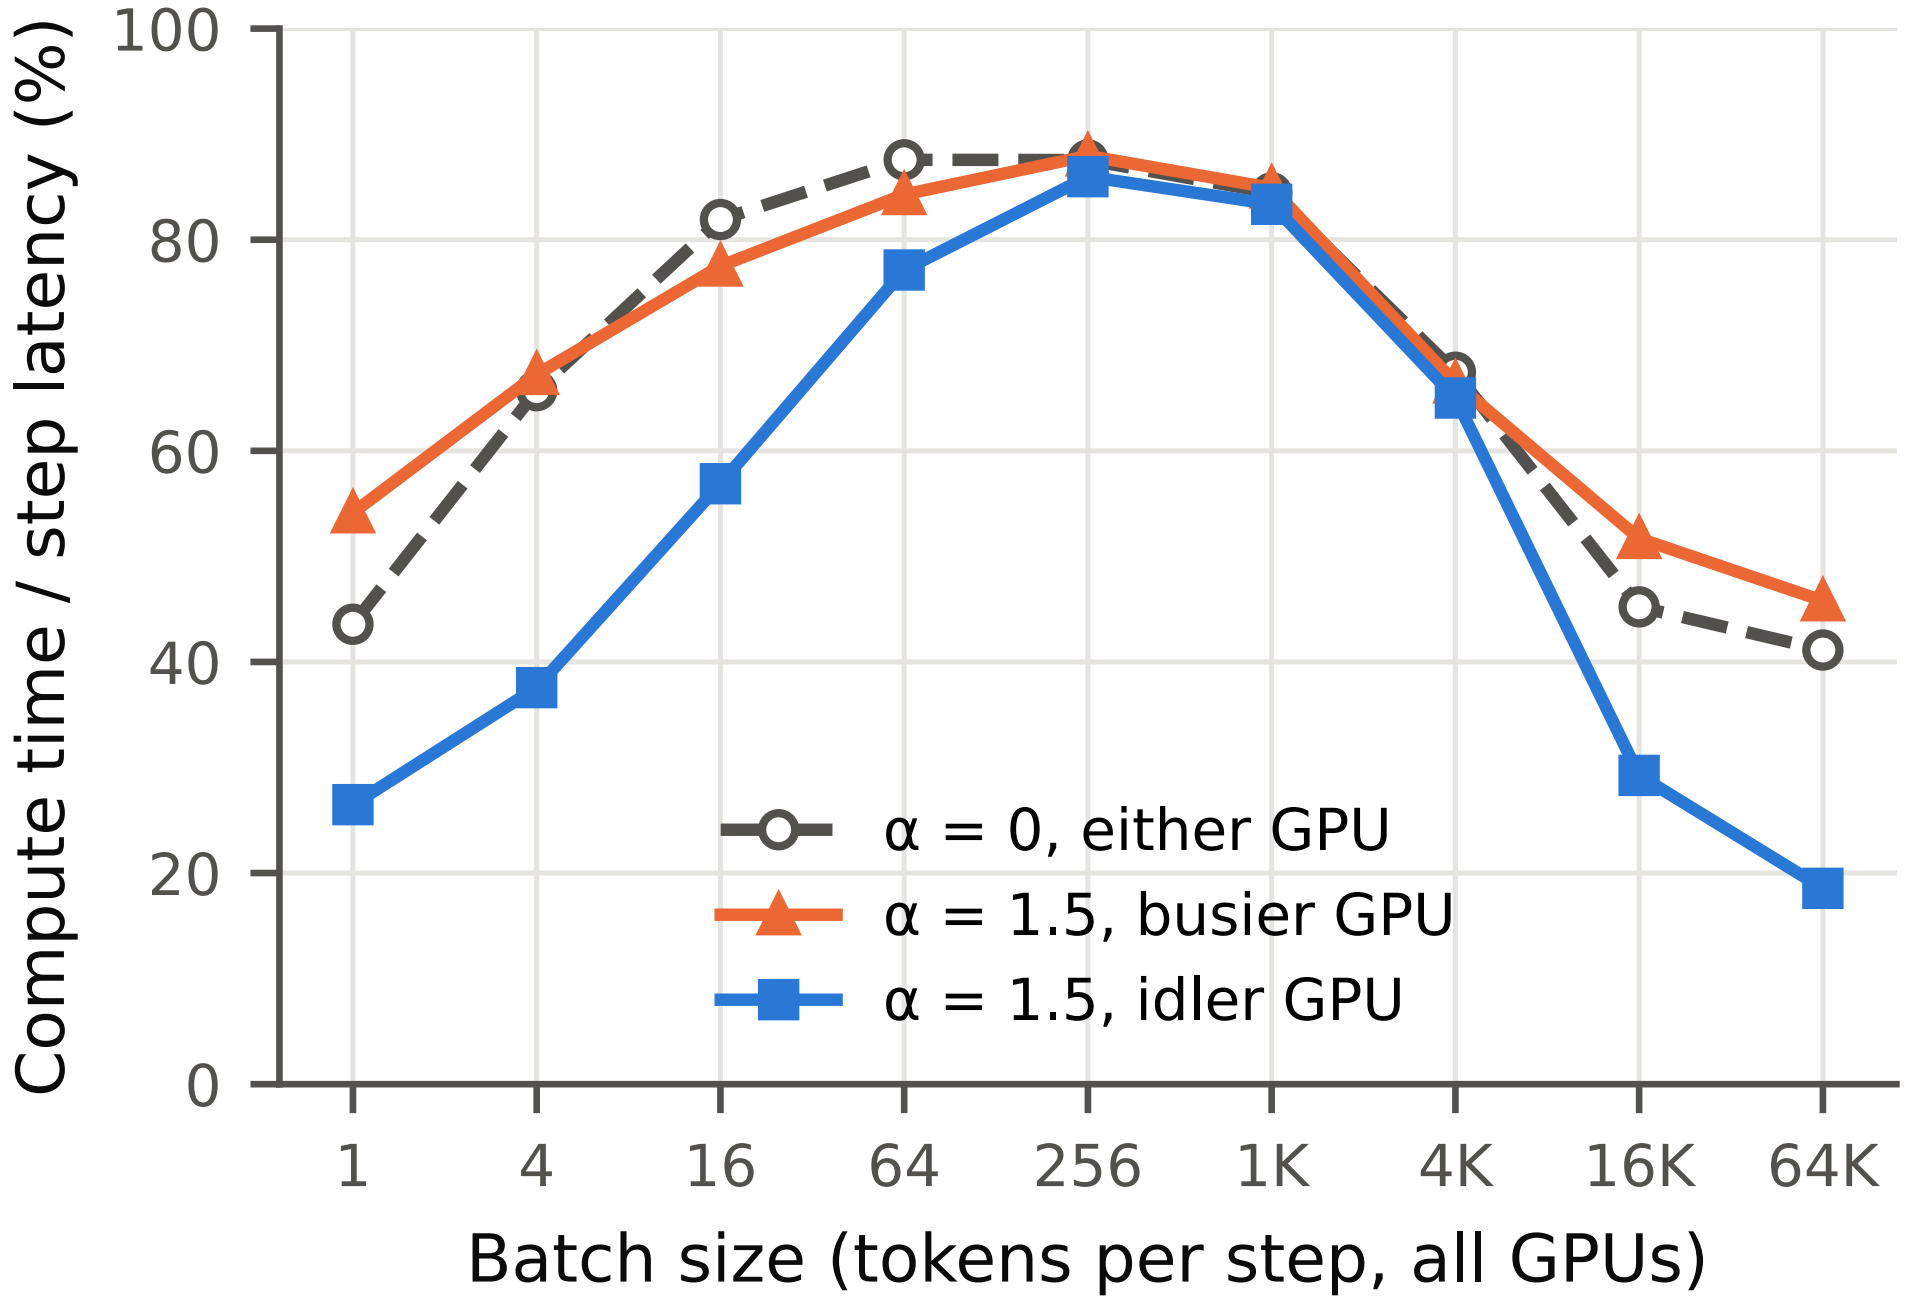
Compute time / step latency (%)
100
80
60
40
20
0
1
4
16
64
256
1K
4K
16K
64K
Batch size (tokens per step, all GPUs)
α = 0, either GPU
α = 1.5, busier GPU
α = 1.5, idler GPU