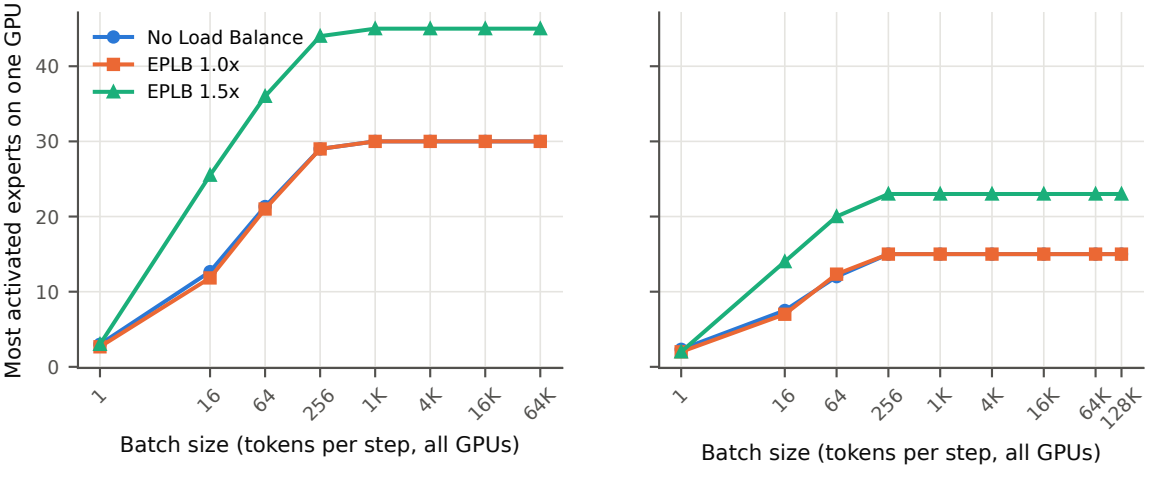
(a) 2 GPUs, α = 1.5
(b) 4 GPUs, α = 1.5
No Load Balance
EPLB 1.0x
EPLB 1.5x
Most activated experts on one GPU
0
10
20
30
40
1
16
64
256
1K
4K
16K
64K
128K
Batch size (tokens per step, all GPUs)
Batch size (tokens per step, all GPUs)