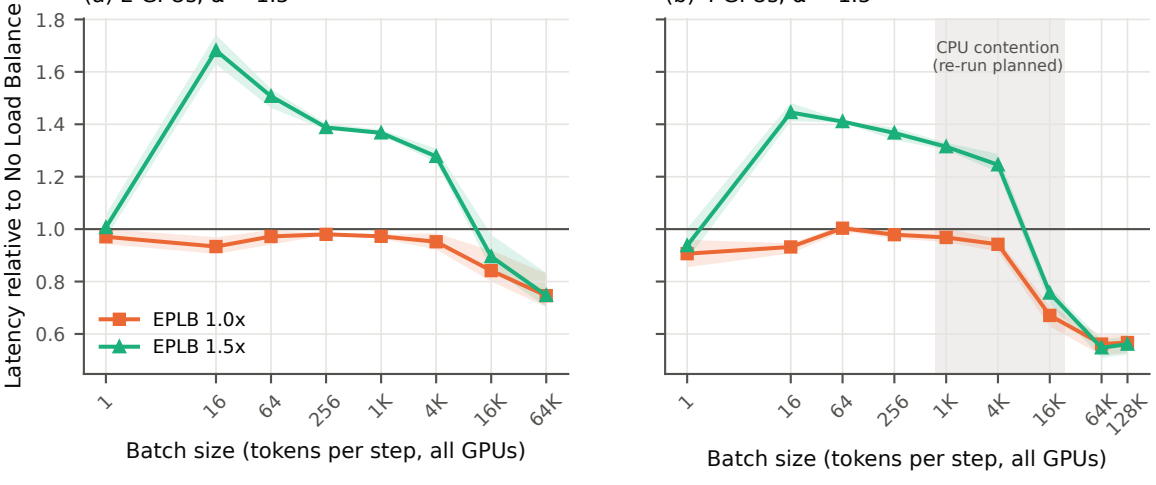
(a) 2 GPUs, α = 1.5
(b) 4 GPUs, α = 1.5
Latency relative to No Load Balance
1.8
1.6
1.4
1.2
1.0
0.8
0.6
EPLB 1.0x
EPLB 1.5x
CPU contention
(re-run planned)
1
16
64
256
1K
4K
16K
64K
128K
Batch size (tokens per step, all GPUs)
Batch size (tokens per step, all GPUs)