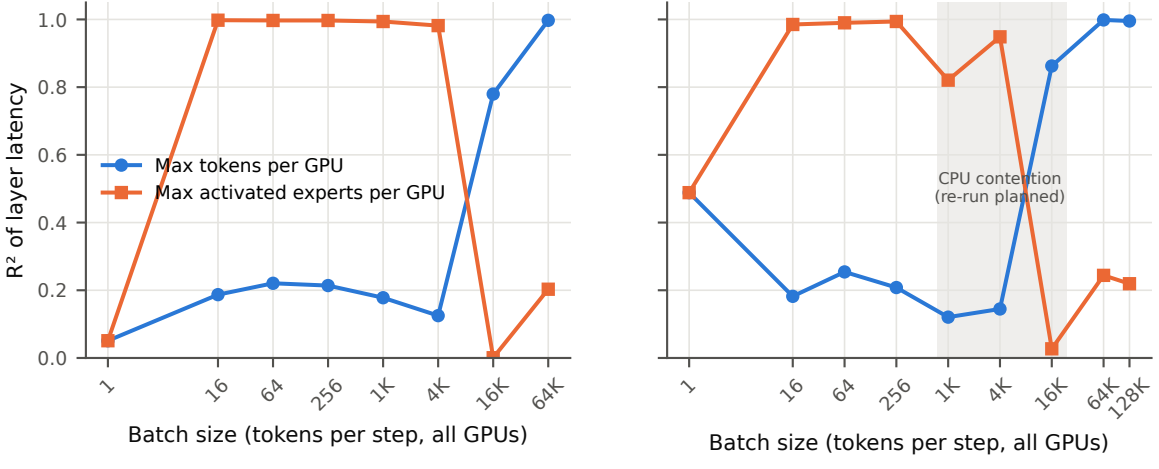
(a) 2 GPUs
(b) 4 GPUs
R² of layer latency
1.0
0.8
0.6
0.4
0.2
0.0
Max tokens per GPU
Max activated experts per GPU
CPU contention
(re-run planned)
1
16
64
256
1K
4K
16K
64K
128K
Batch size (tokens per step, all GPUs)
Batch size (tokens per step, all GPUs)